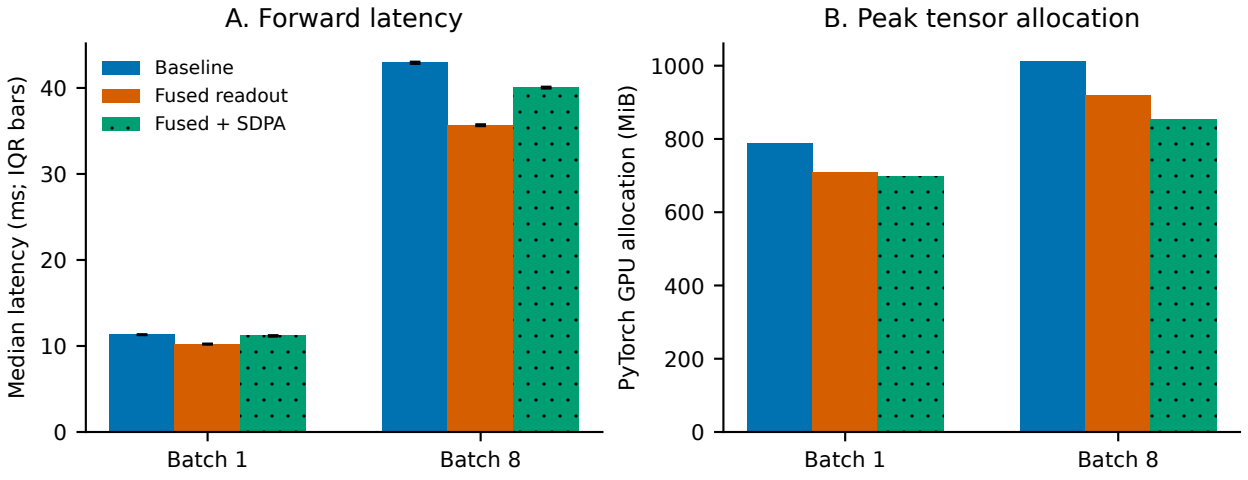
A. Forward latency
Baseline
Fused readout
Fused + SDPA
Median latency (ms; IQR bars)
0
10
20
30
40
Batch 1
Batch 8
B. Peak tensor allocation
PyTorch GPU allocation (MiB)
0
200
400
600
800
1000
Batch 1
Batch 8